

Software Craftsmanship for the Curious Craftsperson

David Souther

April 29, 2012

Contents

0	Introduction	1
0.1	Computers Are Tools	2
0.1.1	How to use this tool	2
0.1.2	How to share this tool	3
0.2	Computer Languages	3
0.2.1	C	4
0.2.2	Python	4
0.2.3	Javascript	5
0.3	Overview of the Text	5
0.3.1	Syntactic Core (1–4)	5
0.3.2	Programming Patterns (5–9)	5
0.3.3	Building Large Programs (10–13)	5
0.3.4	Appendices: Tools for the Toolmakers	6
0.4	Workbooks	6
I	Syntactic Core	7
1	Basic Types and Control Flow	9
1.1	Syntax	9
1.1.1	Keywords	10
1.1.2	Variables	10
1.1.3	Operators	10
1.2	Basic Types	10
1.2.1	Numbers	11
1.2.2	Characters	12
1.2.3	Strings	13
1.3	Control Flow	13
1.3.1	Blocks	13
1.3.2	Branching	13
1.3.3	Looping	13
1.4	Comments	13

2	Functions, Arrays, and Strings	15
2.1	Functions	15
2.1.1	Procedure	15
2.1.2	Function	15
2.1.3	Recursion	15
2.2	Arrays	15
2.3	Strings	15
3	Source Control	17
3.1	Tracking	17
3.2	Sharing	17
3.3	Fixing Goofs	17
4	Introduction to Objects	19
4.1	Structured Data	19
4.2	Methods	19
II	Programming Patterns	21
5	Object Composition	23
5.1	Inheritance	23
5.2	Polymorphism	23
5.3	Visibility	23
6	Debugging	25
6.1	Logging	25
6.2	Exceptions	25
6.3	Breakpoints	25
6.3.1	Watches	25
7	Input & Output	27
7.1	Input	27
7.2	Output	27
7.3	Parameters	27
7.4	Command Line	27
7.5	Files	27
7.6	Databases	27
8	Patterns I: Lazy Loading, Singleton, Iterators	29
8.1	Lazy Loading	29
8.2	Singleton	29
8.3	Iterators	29

9 Patterns II: Factory, Decorator, Visitor	31
9.1 Factory	31
9.2 Decorator	31
9.3 Visitor	31
III Building Large Programs	33
10 Unit Testing	35
11 Model-View-Controller	37
12 MVC Redux	39
13 Team Programming	41
IV Appendices: Tools for the Toolmakers	43
A Text Editors	45
A.1 Vi	45
A.2 gedit	45
B Linux	47
B.1 Linux Desktop	47
B.2 Command Line	47
B.2.1 Vi	47
B.2.2 Bash	47
C Computer Hardware	49
C.1 Computer Arithmetic	49
C.2 Von-Neuman Machine	49
C.3 High-Level Assembly	49
C.4 Memory Hierarchy	49

Chapter 0

Introduction

Human endeavors rest on the backs of the hard-working crafters. From stone-movers of the Egyptian pyramids to steel-workers on today's oil rigs, skilled workers built civilization with their hands. Stone, steel, lumber, and leather have for centuries been the foundation of human enterprise. In the 21st century, there is a new medium demanding attention: information. The data flying across the Internet is the backbone of international trade and commerce, and needs skilled craftspeople to shape it. Yet even as professional carpenters build masterwork cabinetry for law firms and movie stars, there are laymen working on the same craft with the same tools in their garage. This is no less true for computers — the relatively low cost of consumer software gives hobbyist programmers the same tools to work with as the professional.

This book is aimed at those who are interested in this new medium as a potential hobby or curiosity. The book begins assuming the reader knows how to turn on and use their computer for basic tasks — email, word processing, image editing, and video games, and takes them to a level where they will be comfortable and confident in using and controlling their computers. This book is very fast paced. Software development is a field which has undergone active development for the past 80 years, yet did not exist beyond a dream before then. There is a lot to learn about this craft, but readers who want to become truly skilled at this craft will hopefully find this book gives them many of the tools they need to feel comfortable working their computer. To get the most out of the text, I recommended readers work through the companion workbook. Like any craft, to get good at software development you need to develop software. The workbooks lead you through developing software, in a way highlights the concepts presented in the text.

If the book seems hard, fear not. Rome wasn't built in a day, and research shows it takes ten years of practicing a skill to truly develop expertise. This book is one way to approach learning this subject, which hopefully will work for you. If not, one of the many references in the bibliography might provide a different approach that works better for you.

0.1 Computers Are Tools

Information is a critical piece of today's infrastructure. How people use information is part of the field of Computer Science. Computer Science is on the one hand deeply seated in mathematics, and on the other, firmly rooted in practicality. The computer itself is simply a tool, like a band saw or plasma torch, which do amazing things with this data. Whether used in large-scale data mining and predictions for financial institutions, determining exactly what websites have what content, or creating a side show of family photos. At the end of the day, the laptop or desktop sitting in front of you is just a piece of silicone, copper, and plastic, capable of doing only exactly what it is told.

The task of the computer programmer is to tell the computer what to do. This is not an easy task. When people think about a problem, we can start off being a bit loose on how we describe the problem to ourselves. We can find issues or errors in our assumptions, and change them on the fly. Computers cannot do this. Every condition must be considered before hand. Say you're balancing your check book. You add up \$12.57, \$52.45, and \$1.99, but miss the decimal on the \$52.45, and end up with \$5249.56 — clearly an error. On paper, it is obvious what you did wrong. The computer has no way of knowing this was an error. Instead, it is up to the good programmer to tell the computer to verify that any numbers typed into the financial program ends with a decimal and two numbers. Then, the program can warn the user before making the calculation, potentially avoiding a costly transaction.

0.1.1 How to use this tool

Software craftsman use this tool by writing programs. A program is a document both written and read by human beings, while telling the computer unambiguously what to do. Take this example: $4 + 6/2$. Is the answer 7 or 5? If you remembered something like "Please Remember My Dear Aunt Sally" from grade school, you would say 7. If you were a desk calculator in an accountant's office, you would say 5. This statement is *ambiguous* — it could mean more than one thing. The programmer's job is to decide if the statement should be $(4 + 6)/2 = 5$ or $4 + (6/2) = 7$.

Of course, this is a minor pedantic exercise. As software gets more complex, the question becomes "how does the computer help people interact with their data?" Facebook is an excellent tool for managing data regarding many aspects of your relationships with friends and relatives online, by making it very easy to share photos and other tidbits of information. Microsoft Word is an excellent program to enable people to write documents, vital pieces of data for business, education, and private correspondence. Clearly, there is much to consider when writing a program — when using the tools of software craftsmanship.

0.1.2 How to share this tool

There is a whole world outside your front doorstep, and as much as some programmers might want to deny it, at some point every craftsman's programming will be used by someone else. It is important your code be both readable and usable by these other craftsmen. There are many ways to do this, and many great tools to facilitate this. While many software craftsman may want to hoard their code, this is simply not possible. Even at companies like Microsoft, the largest proprietary software development company on the planet, teams are used because software projects, even most hobbyist projects, are too large for one person to handle alone.

Many computer science courses ignore this aspect of programming, which I believe is a real detriment to the students. Many computer science students can go their entire undergraduate careers without ever looking at their peer's code, and only having their program reviewed by the instructors. Craftsman of all fields can drastically improve their work by meeting and sharing ideas. Building partnerships with other craftsmen is one of the most rewarding ways to practice a trade, but can also be a harrowing experience learning how to work and interact with a completely new group of people. To work past this impediment, I present tools for managing and sharing your programs early in the book (Chapter 3, Source Control), and encourage you to find other programmers and technology clubs in your area and on the Internet.

0.2 Computer Languages

Computer languages are the tie between the world of everyday common sense conversation and the exact stupidity of the computer. There are hundreds, if not thousands, of different programming languages today. Many are “toy” languages, built for a fun exercise or class project and are not intended for wide-scale use. There are other languages that are proven as the working horses of computers, and are used everywhere from your cell phone to robots running on Mars.

Languages are often classed based on their style of programming, and their level of expressivity. There are, broadly, three styles of programming: imperative, functional, and logical. Imperative languages are similar to a cookbook recipe. They describe, one statement at a time, the “things” a computer is to do—add two numbers, print those numbers to the screen, ask the user for confirmation. C, JavaScript, and Python are all imperative languages. Functional languages embrace the mathematics of computer science. Functional languages often use similar syntactic constructs as imperative languages, but with fundamental underlying differences. Haskell is a functional programming language. Logical programming languages have little to do with either imperative or functional programming. Prolog is a logical programming language. Logical programming will not be discussed in this book.

Expressivity is notion for the ratio of amount of programming code written

to how much the computer does. Languages with the least amount of expressivity are referred to as “low-level” languages — they must deal with all aspects of the computer and its memory. “High-level” languages handle many of the details of working on computer hardware, and let the programmer just focus on expressing the logic of the program in question. This book is presented using three different languages. These languages were chosen because they each represent a very different approach to software craftsmanship. Each approach is correct in its own way, and it is important to know each of the approaches to be a great software craftsman. Further, each of these three languages is mature, and actively used in a variety of projects today.

0.2.1 C

C started its life at Bell Labs in 1973. C was written by Dennis Kernighan and Brian Ritchie as a language to write the Unix operating system in. Before C, operating systems (the program enabling all the other programs to run) were written in an assembly language for each computer Unix would run on. However, no two brands of computers had the same assembly language, so any time someone wanted to run Unix on a new computer, they had to rewrite the entire operating system. At the time, this could be thousands of lines of code. Today, it would take millions of lines of assembly to code Windows or Linux. C was designed to be a consistent language that, rather than being run on a computer as is, would first be compiled into the appropriate assembly code.

C is widely considered the “lowest-level” of today’s common programming languages, meaning C is as close to running “on the hardware” as you can get. When writing C, the programmer has to deal with many aspects of computer memory management. There are few utilities to achieve all but the most common tasks (though there are a wealth of “libraries” to fill the gap). C is the only language used in this book that must be compiled before being run. This “low-level” nature of C makes it a very powerful language, especially when faced with requirements to interact directly with hardware, or when hardware is in short supply (embedded on robots or cell phones). That power comes with great responsibility for writing the program correctly.

0.2.2 Python

Python is a programming language developed by Guido van Rossum in the late 1980s. Python has gone through two major revisions since its first release, and now is widely available on nearly any computing platform. Today, Python is used by many Linux distributions to write a variety of their higher-level tools. Organizations from Google to NASA use Python for a variety of their mission-critical applications.

Python was designed to be a very flexible language, and as such is high-level compared to C. Python was also designed to be a fun language to use—the name Python refers not to the snake, but to Monty Python’s Flying Circus. Python has a certain culture around its use not seen in many other programming

languages. In Python, the prevailing wisdom is “there should be one—and preferable only one—obvious way to do it.”

0.2.3 Javascript

Javascript is a programming language developed at Netscape in 1994 by Brandon Eiche over the course of a week. Javascript is the de-facto standard for writing programs served over the Internet to run in a user’s web browser. While actually a functional programming language, Javascript is often (mis)represented as being an imperative language. This has lead to some very poor code being written in the 15 plus years since its inception. That said, its use in every web browser today has many people working to make Javascript a less-maligned and better respected language.

Javascript is a very high level language. The programmer has few worries about memory management, and no capabilities to access the computer’s hardware (though there are initiatives to enable such use). When combined with libraries like jQuery, javascript can be the most expressive of the three languages presented. When we start working with graphical programs later in the book, Javascript’s expressive power will really shine, in that the amount of code needed to do the same thing (click a button) can often be an order of magnitude less than a similar program in C.

0.3 Overview of the Text

0.3.1 Syntactic Core (1–4)

The first four chapters will cover the core of writing programs. This will get readers to the point where they can have their computer talking to them and asking questions, though perhaps not gracefully. This section also covers the basics of source control, and sharing code between developers.

0.3.2 Programming Patterns (5–9)

The second section covers patterns in software craftsmanship—commonly used ideas which can be the large building blocks of a program’s design. This includes discussions on modeling data, saving data, and some ideas on how to make good program design decisions. This section also includes a chapter on debugging—finding out what went wrong with a program (things will go wrong).

0.3.3 Building Large Programs (10–13)

The last section deals with managing large pieces of software. The section will go through building a basic “paint” application, where users can draw lines and colors on a canvas, and save the painting to look at later. This section also includes a chapter on how to work with other people on the same project.

0.3.4 Appendices: Tools for the Toolmakers

Any craftsman needs a workbench that suits their needs. The appendices address several aspects around such a workbench. First and foremost, programs are text. As such, it is important to use a text editor that facilitates looking at and ‘seeing’ your code. Second is a discussion of the Linux operating system. Its open source nature encourages users to take responsibility for their computer, and provides a wealth of tools for programming unavailable (or at least difficult (read, costly) to acquire) on Windows. Finally, there is a discussion on computer hardware. This is a look at what physically is going on in your program, focusing on the CPU, memory, and some peripherals. This is not strictly critical for learning software programming as a craft, but hopefully you will be interested enough by the end of the book to want to take a deeper look at what’s going on.

0.4 Workbooks

Like any skill, software craftsmanship takes practice. The workbooks are designed to highlight the concepts presented in the text, while giving you an opportunity to practice these skills. The workbooks are broken into lessons roughly corresponding to sections of the main text. Each lesson has two parts. The first part is a listing of code. You should type the code into your editor exactly as written. Do not copy and paste the code. Much of programming involves paying very close attention to a myriad of small details, and every character has meaning. This discipline in typing code exactly as presented will pay off in your programming future. The second part of each lesson is a few exercises to work with the new concepts in the lesson, and combine them with what you learned and wrote previously. Some of the lesson exercises will involve conducting research on the internet. Being able to find help with programming questions is another invaluable skill as a software developer.

Part I

Syntactic Core

Chapter 1

Basic Types and Control Flow

Computers work with data. At the lowest level, this data is just ‘on’ and ‘off’ in a transistor. Software craftsmen don’t work at this level. Instead, our programming languages give us the tools to work with this data in a much more intuitive fashion. In broad strokes, this is achieved with two concepts: data types and control flow. Data types describe the data our program is working with, whether it be a number, a bank account, or an image. Control flow describes the operations that occur on the data over time, like adding numbers, checking the balance of a bank account, or cropping an image.

You can think of data types as the space of a program. At any one point in time, we will have certain sets of data we are working with. We can stop the program, look at the data, and carry on. While we think of three dimensions in space, up/down, left/right, and front/back, there are many more dimensions we could talk about for different things. A table has a color, along with its position. Our programs represent these properties in different ways, but we as programmers abstract them, and the term we use to talk about these properties is data types.

Control flow would then be the time aspect of a program. As the program runs, it does different things to the data. At some point, it might need to make a decision, and do one thing or another depending on what properties the data has. At another point, the program might need to do the same thing on a bunch of similar pieces of data. Control flow represents the logical actions our programs follow.

1.1 Syntax

Before we dive into discussing our first program, there is one more topic to cover. Syntax describes the words and symbols that make a valid program. In English, nouns, verbs and adjectives must come in a certain fashion for a

sentence to make sense. In a computer program, there are similar rules for how a program can be written, and still make sense to the computer to run. We will cover the details of syntax for each programming language in the workbook, but there are some key terms to know regardless of the language.

1.1.1 Keywords

In any programming language, there are a few reserved *keywords* — words which have a specific meaning to the program. Most programming languages only use no more than a couple dozen keywords. These will be covered in the workbooks as needed.

1.1.2 Variables

In a program file, variables are words used to refer to some piece of data. Variables let programmers use concrete concepts to refer to the more abstract realm of the mathematical values in the computer. Variables can generally be any word that isn't a keyword in a programming language.

1.1.3 Operators

Operators are the built-in things that can happen to data. Common operators are things from elementary arithmetic — add +, subtract −, multiply *, and divide /. Other operators let the program compare two pieces of data — equals (==), less-than (<), greater-than- or- equal-to (>=). Notice there are some funny things for the comparison operators. First, equals is two characters, ==, not the single = used in elementary school. This is because a single = sign is an operator used to assign values to a variable. Multiplication is a star, *. While x was used in elementary school, we use * for our programs because x could be a variable, and we need to be unambiguous in what the program means. There is no ≤ on a common keyboard, so programs use the two character combinations of <= and >= to do comparisons. There are other operators that will be covered as they are needed.

1.2 Basic Types

Data Types fall into two distinct categories. ‘Primitive’ data types are those that are specified in the programming language itself. ‘Complex’ data types are defined by the programmer, and are made up entirely of primitive data types and other complex data types already defined. In this section, we will go over the basic data types in your language. In general, these basic types are applicable in any programming language in major use today, with small discrepancies between them.

1.2.1 Numbers

Since computers only do arithmetic, everything is at some level a number in the computer. In mathematics, numbers can be as big as they need to be. If we were counting pebbles of sand on a beach, there is a number to represent that (though I will not be the one doing the counting). Computers are somewhat more finite. The representation of a basic number in a computer is confined to a range of a few possible values. On today's computers, that range is either 2^{32} or 2^{64} possible values. The 32 and 64 refer to the number of 'bits' your computer handles. When the IT guy asks whether your computer is 32 or 64 bit, this is what we're talking about.

Integers

Integers in computer programming are Whole Numbers from elementary school — 1, 5, -15, 47. They can be positive or negative. In some languages, they are limited in size based on your computer's processor. On a 32 bit computer, the range of integers is from -2^{31} (-2 147 483 648) to $2^{31} - 1$ (2 147 483 647) (the one positive number is lost for 0). On a 64 bit computer, the range of values is -2^{63} to $2^{63} - 1$. It is left to the reader to determine those numbers in base 10. Integers are useful in many cases because they are very easy to work with, but have a major drawback of being unable to work with fractional parts. That is the use of Floating-point numbers.

Floats

In your computer, there is only so much room to store numbers. As mentioned in the previous section on integers, there are only 2^{32} possible numbers you can represent on a 32-bit computer. Those integers are only whole numbers, and can't represent a number like $\frac{1}{2}$ or $\pi = 3.1415926\dots$. Instead, we use a special type of number called a *Floating Point* number. A floating point number can still only represent 2^{32} different values, but it is interpreted in such a way that it can represent fractional numbers. At this point, we must take a detour into the world of mathematics in order to understand how floating point numbers work, and in the process save ourselves from making some very costly mistakes.

Floating Point Arithmetic This section explains some complicated math.

Our everyday number system is a base-10 number system. When you write your numbers, you write 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12 Notice how starting at 10 we use two characters to represent the number- 1 and 0. "Well, duh" you might say. In fact, there is a very specific reason we need two numbers at this point. What 10 means is $(1 * 10^1) + (0 * 10^0)$. This is the base-10 system- each character in the number means $a * 10^p$ where a is the number, and p is its position (starting from 0). So, the number 347 is $(3 * 10^2) + (4 * 10^1) + (7 * 10^0)$. In a computer, we use binary, or a base-2 number system. Instead of writing numbers as $a * 10^p$, we write them as $a * 2^p$. So, to get 12, we need to write 1100 , or $(1 * 2^3) + (1 * 2^2) + (0 * 2^1) + (0 * 2^0) = (1 * 8) + (1 * 4) + (0 * 2) + (0 * 1) = 8 + 4 + 0 + 0 = 12_{10}$.

We can talk about decimal numbers like 1.56 in the same way— $(1 * 10^0) + (5 * 10^{-1}) + (6 * 10^{-2}) = 1 + \frac{5}{10} + \frac{6}{100} = \frac{100}{100} + \frac{50}{100} + \frac{6}{100} = \frac{156}{100} = \frac{39}{25}$.

While this math might be disheartening, it is important to remember that computers do work in this fashion under the hood. That said, this is the only time this will be discussed in the main book (though some advanced exercises in the workbook will deal with computer arithmetic in more depth, and appendix 3 is devoted to a discussion of just how the computer does work under the hood).

1.2.2 Characters

Characters are single bytes of data that are interpreted as being human-readable in some form. For instance, the letter ‘a’ is the integer 97. The letter ‘V’ is 86. ‘!’, the exclamation point, is 33. Because the computer operates on numbers, there is no distinction to it between the number 33 and the character ‘!’. Instead, programmers tell their programs how they intend to use the data, and the program does the “right thing”. The “right thing” is exactly what is was told to do. To understand what you’re telling the computer to do better, it is important to know that there is nothing special about the mappings between the numbers and characters. Indeed, the first 128 character numbers were chosen by an organization called the American Standard Code for Information Interchange, or ASCII. For this reason, the first 128 characters are called the “ASCII” character set. Of course, there is no way to fit all the characters of both the Latin and Cyrillic alphabets into 128 different numbers, much less most Asian character sets. For this reason, there is another character set called UNICODE that defines over 4 million possible characters, enough to satisfy human languages for some time. We will not go into the details of Unicode here.

ASCII Code Chart

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	NUL	SOH	STX	ETX	EOT	ENQ	ACK	BEL	BS	HT	LF	VT	FF	CR	SO	SI
1	DLE	DC1	DC2	DC3	DC4	NAK	SYN	ETB	CAN	EM	SUB	ESC	FS	GS	RS	US
2		!	"	#	\$	%	&	'	()	*	+	,	-	.	/
3	0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4	@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5	P	Q	R	S	T	U	V	W	X	Y	Z	[\]	^	_
6	`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7	p	q	r	s	t	u	v	w	x	y	z	{		}	~	DEL

Let’s take a closer look at some features of the ASCII character set. First, notice that uppercase and lowercase characters are represented distinctly. This is why case sensitivity is important on computers. Second, 0 through 31 look really funky. That’s because these are called control characters, not printing characters. Most of these control characters are no longer used in today’s computers (they were used to literally control printers and other devices in the 70s and 80s), but there are some that warrant special attention. The first is NULL,

0. NULL is used in C to represent the end of an array (see Chapter 2). 10, ‘\n’ or Newline, is used to make the computer put a line break in a program’s output; otherwise, everything would show up on one long horizontal line. Second is 15, ‘\r’ or Carriage Return. This goes back to when computer printers were fancy typewriters, which required a special symbol to move the carriage back to the starting position. While I can’t think of any printers today that need this functionality, it is not uncommon to see ‘\r\n’ as a legacy pair of characters in many programs. The last control character that is interesting to us is 9, ‘\t’, horizontal tab. The horizontal tab tells the computer that we want a lot of whitespace (usually 4 or 8 spaces worth).

1.2.3 Strings

Characters on their own aren’t particularly interesting. What is useful is having long runs of characters to make up words, sentences, paragraphs, and so on. We achieve this by using what are called strings— which are literally just a collection of characters in a row. Usually strings are surrounded by double quotes in a program. See the workbook for a variety of examples on how strings look and work in a program.

While characters really are treated as just numbers, strings have a variety of common operations that don’t make sense with numbers. Among these are operations to determine the length of a string (how many characters it has), to combine strings together, and to pull strings apart. These are also covered in detail in the workbook.

1.3 Control Flow

1.3.1 Blocks

Expression

Statement

1.3.2 Branching

If-Then-Else

1.3.3 Looping

For Range

While

1.4 Comments

There is one last piece of a program we need to mention. Comments are text in our programs that are not used by the computer, but are purely for

Chapter 2

Functions, Arrays, and Strings

2.1 Functions

2.1.1 Procedure

2.1.2 Function

2.1.3 Recursion

2.2 Arrays

2.3 Strings

Chapter 3

Source Control

3.1 Tracking

3.2 Sharing

3.3 Fixing Goofs

Chapter 4

Introduction to Objects

4.1 Structured Data

4.2 Methods

Part II

Programming Patterns

Chapter 5

Object Composition

5.1 Inheritance

5.2 Polymorphism

5.3 Visibility

Chapter 6

Debugging

6.1 Logging

6.2 Exceptions

6.3 Breakpoints

6.3.1 Watches

Chapter 7

Input & Output

7.1 Input

7.2 Output

7.3 Parameters

7.4 Command Line

7.5 Files

7.6 Databases

Chapter 8

Patterns I: Lazy Loading, Singleton, Iterators

8.1 Lazy Loading

8.2 Singleton

8.3 Iterators

Chapter 9

Patterns II: Factory, Decorator, Visitor

9.1 Factory

9.2 Decorator

9.3 Visitor

Part III

Building Large Programs

Chapter 10

Unit Testing

Chapter 11

Model-View-Controller

Chapter 12

MVC Redux

Chapter 13

Team Programming

Part IV

Appendices: Tools for the Toolmakers

Appendix A

Text Editors

A.1 Vi

A.2 gedit

Appendix B

Linux

B.1 Linux Desktop

B.2 Command Line

B.2.1 Vi

B.2.2 Bash

Appendix C

Computer Hardware

C.1 Computer Arithmetic

C.2 Von-Neuman Machine

C.3 High-Level Assembly

C.4 Memory Hierarchy